

```
2017-02-11 01:30:41,066 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to
ResourceManager at ip-172-31-56-198.ec2.internal/172.31.56.198:8032
2017-02-11 01:30:44,776 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input paths to process : 1
2017-02-11 01:30:44,788 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2017-02-11 01:30:44,792 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev f7cb0596948c5bfd3e71d37b0f5bb21a19554666]
2017-02-11 01:30:45,233 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:17
2017-02-11 01:30:45,509 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1486776510473_0001
2017-02-11 01:30:46,383 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1486776510473_0001
2017-02-11 01:30:46,517 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-56-198.ec2.internal:20888/proxy/application_1486776510473_0001/
2017-02-11 01:30:46,518 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1486776510473_0001
2017-02-11 01:30:58,646 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776510473_0001 running in uber mode : false
2017-02-11 01:30:58,647 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2017-02-11 01:31:17,780 INFO org.apache.hadoop.mapreduce.Job (main): map 4% reduce 0%
2017-02-11 01:31:18,784 INFO org.apache.hadoop.mapreduce.Job (main): map 5% reduce 0%
2017-02-11 01:31:21,797 INFO org.apache.hadoop.mapreduce.Job (main): map 7% reduce 0%
2017-02-11 01:31:22,812 INFO org.apache.hadoop.mapreduce.Job (main): map 9% reduce 0%
2017-02-11 01:31:24,820 INFO org.apache.hadoop.mapreduce.Job (main): map 14% reduce 0%
2017-02-11 01:31:25,823 INFO org.apache.hadoop.mapreduce.Job (main): map 16% reduce 0%
2017-02-11 01:31:26,827 INFO org.apache.hadoop.mapreduce.Job (main): map 19% reduce 0%
2017-02-11 01:31:27,832 INFO org.apache.hadoop.mapreduce.Job (main): map 23% reduce 0%
2017-02-11 01:31:28,837 INFO org.apache.hadoop.mapreduce.Job (main): map 29% reduce 0%
2017-02-11 01:31:29,842 INFO org.apache.hadoop.mapreduce.Job (main): map 36% reduce 0%
2017-02-11 01:31:30,846 INFO org.apache.hadoop.mapreduce.Job (main): map 39% reduce 0%
2017-02-11 01:31:31,849 INFO org.apache.hadoop.mapreduce.Job (main): map 43% reduce 0%
2017-02-11 01:31:32,853 INFO org.apache.hadoop.mapreduce.Job (main): map 50% reduce 0%
2017-02-11 01:31:33,857 INFO org.apache.hadoop.mapreduce.Job (main): map 52% reduce 0%
2017-02-11 01:31:34,861 INFO org.apache.hadoop.mapreduce.Job (main): map 54% reduce 0%
2017-02-11 01:31:35,868 INFO org.apache.hadoop.mapreduce.Job (main): map 60% reduce 0%
2017-02-11 01:31:36,875 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 0%
2017-02-11 01:31:37,879 INFO org.apache.hadoop.mapreduce.Job (main): map 66% reduce 0%
2017-02-11 01:31:38,883 INFO org.apache.hadoop.mapreduce.Job (main): map 70% reduce 0%
2017-02-11 01:31:39,887 INFO org.apache.hadoop.mapreduce.Job (main): map 74% reduce 0%
2017-02-11 01:31:40,896 INFO org.apache.hadoop.mapreduce.Job (main): map 77% reduce 0%
2017-02-11 01:31:41,900 INFO org.apache.hadoop.mapreduce.Job (main): map 78% reduce 0%
2017-02-11 01:31:43,907 INFO org.apache.hadoop.mapreduce.Job (main): map 80% reduce 0%
2017-02-11 01:31:46,919 INFO org.apache.hadoop.mapreduce.Job (main): map 82% reduce 0%
2017-02-11 01:31:47,924 INFO org.apache.hadoop.mapreduce.Job (main): map 84% reduce 0%
2017-02-11 01:31:48,927 INFO org.apache.hadoop.mapreduce.Job (main): map 84% reduce 2%
2017-02-11 01:31:49,931 INFO org.apache.hadoop.mapreduce.Job (main): map 88% reduce 2%
2017-02-11 01:31:50,936 INFO org.apache.hadoop.mapreduce.Job (main): map 94% reduce 2%
2017-02-11 01:31:51,939 INFO org.apache.hadoop.mapreduce.Job (main): map 96% reduce 3%
2017-02-11 01:31:52,943 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 7%
2017-02-11 01:31:53,949 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 7%
2017-02-11 01:31:54,952 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 11%
2017-02-11 01:31:55,955 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 26%
2017-02-11 01:31:56,960 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 33%
2017-02-11 01:31:57,966 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 44%
2017-02-11 01:32:00,976 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 56%
2017-02-11 01:32:03,989 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 78%
2017-02-11 01:32:04,992 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 89%
2017-02-11 01:32:06,049 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-02-11 01:32:07,058 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776510473_0001 completed successfully
2017-02-11 01:32:07,171 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
File System Counters
```

FILE: Number of bytes read=5139602
FILE: Number of bytes written=17158721
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=1683
HDFS: Number of bytes written=0
HDFS: Number of read operations=17
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=1087569576
S3: Number of bytes written=975899
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0

Job Counters

Killed map tasks=2
Killed reduce tasks=1
Launched map tasks=18
Launched reduce tasks=9
Data-local map tasks=18
Total time spent by all maps in occupied slots (ms)=35457792
Total time spent by all reduces in occupied slots (ms)=11628384
Total time spent by all map tasks (ms)=738704
Total time spent by all reduce tasks (ms)=121129
Total vcore-milliseconds taken by all map tasks=738704
Total vcore-milliseconds taken by all reduce tasks=121129
Total megabyte-milliseconds taken by all map tasks=1134649344
Total megabyte-milliseconds taken by all reduce tasks=372108288

Map-Reduce Framework

Map input records=30868726
Map output records=30868726
Map output bytes=895193054
Map output materialized bytes=8718061
Input split bytes=1683
Combine input records=30868726
Combine output records=937202
Reduce input groups=28980
Reduce shuffle bytes=8718061
Reduce input records=937202
Reduce output records=28980
Spilled Records=1874404
Shuffled Maps =153
Failed Shuffles=0
Merged Map outputs=153
GC time elapsed (ms)=14376
CPU time spent (ms)=246350
Physical memory (bytes) snapshot=15876149248
Virtual memory (bytes) snapshot=97854750720
Total committed heap usage (bytes)=14437318656

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1087569576

File Output Format Counters

Bytes Written=975899