

```
2017-02-11 01:29:42,887 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-55-89.ec2.internal/172.31.55.89:8032
2017-02-11 01:29:46,084 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input paths to process : 1
2017-02-11 01:29:46,113 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2017-02-11 01:29:46,123 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev f7cb0596948c5bfd3e71d37b0f5bb21a19554666]
2017-02-11 01:29:46,569 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:17
2017-02-11 01:29:46,844 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1486776458734_0001
2017-02-11 01:29:47,512 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1486776458734_0001
2017-02-11 01:29:47,625 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-55-89.ec2.internal:20888/proxy/application_1486776458734_0001/
2017-02-11 01:29:47,626 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1486776458734_0001
2017-02-11 01:29:58,837 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776458734_0001 running in uber mode : false
2017-02-11 01:29:58,838 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2017-02-11 01:30:18,965 INFO org.apache.hadoop.mapreduce.Job (main): map 2% reduce 0%
2017-02-11 01:30:19,971 INFO org.apache.hadoop.mapreduce.Job (main): map 3% reduce 0%
2017-02-11 01:30:21,981 INFO org.apache.hadoop.mapreduce.Job (main): map 4% reduce 0%
2017-02-11 01:30:22,985 INFO org.apache.hadoop.mapreduce.Job (main): map 6% reduce 0%
2017-02-11 01:30:24,996 INFO org.apache.hadoop.mapreduce.Job (main): map 7% reduce 0%
2017-02-11 01:30:26,009 INFO org.apache.hadoop.mapreduce.Job (main): map 15% reduce 0%
2017-02-11 01:30:27,013 INFO org.apache.hadoop.mapreduce.Job (main): map 17% reduce 0%
2017-02-11 01:30:29,030 INFO org.apache.hadoop.mapreduce.Job (main): map 20% reduce 0%
2017-02-11 01:30:31,045 INFO org.apache.hadoop.mapreduce.Job (main): map 26% reduce 0%
2017-02-11 01:30:32,049 INFO org.apache.hadoop.mapreduce.Job (main): map 34% reduce 0%
2017-02-11 01:30:33,054 INFO org.apache.hadoop.mapreduce.Job (main): map 38% reduce 0%
2017-02-11 01:30:34,058 INFO org.apache.hadoop.mapreduce.Job (main): map 41% reduce 0%
2017-02-11 01:30:35,062 INFO org.apache.hadoop.mapreduce.Job (main): map 47% reduce 0%
2017-02-11 01:30:36,066 INFO org.apache.hadoop.mapreduce.Job (main): map 49% reduce 0%
2017-02-11 01:30:37,071 INFO org.apache.hadoop.mapreduce.Job (main): map 55% reduce 0%
2017-02-11 01:30:38,077 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 0%
2017-02-11 01:30:39,081 INFO org.apache.hadoop.mapreduce.Job (main): map 65% reduce 0%
2017-02-11 01:30:40,085 INFO org.apache.hadoop.mapreduce.Job (main): map 71% reduce 0%
2017-02-11 01:30:41,089 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 1%
2017-02-11 01:30:42,094 INFO org.apache.hadoop.mapreduce.Job (main): map 86% reduce 1%
2017-02-11 01:30:43,098 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 1%
2017-02-11 01:30:44,102 INFO org.apache.hadoop.mapreduce.Job (main): map 92% reduce 3%
2017-02-11 01:30:45,106 INFO org.apache.hadoop.mapreduce.Job (main): map 95% reduce 3%
2017-02-11 01:30:46,110 INFO org.apache.hadoop.mapreduce.Job (main): map 97% reduce 3%
2017-02-11 01:30:47,115 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 3%
2017-02-11 01:30:50,125 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 11%
2017-02-11 01:30:52,134 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 22%
2017-02-11 01:30:54,141 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 33%
2017-02-11 01:30:55,145 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 56%
2017-02-11 01:30:57,153 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 78%
2017-02-11 01:30:59,160 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 89%
2017-02-11 01:31:00,365 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-02-11 01:31:00,378 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776458734_0001 completed successfully
2017-02-11 01:31:00,505 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
File System Counters
    FILE: Number of bytes read=13411123
    FILE: Number of bytes written=33630425
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1683
    HDFS: Number of bytes written=0
```

HDFS: Number of read operations=17
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=1087603226
S3: Number of bytes written=957920
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0

Job Counters

Killed map tasks=2
Killed reduce tasks=1
Launched map tasks=18
Launched reduce tasks=9
Data-local map tasks=18
Total time spent by all maps in occupied slots (ms)=30741984
Total time spent by all reduces in occupied slots (ms)=12135744
Total time spent by all map tasks (ms)=640458
Total time spent by all reduce tasks (ms)=126414
Total vcore-milliseconds taken by all map tasks=640458
Total vcore-milliseconds taken by all reduce tasks=126414
Total megabyte-milliseconds taken by all map tasks=983743488
Total megabyte-milliseconds taken by all reduce tasks=388343808

Map-Reduce Framework

Map input records=30868726
Map output records=1647340
Map output bytes=47693990
Map output materialized bytes=16923444
Input split bytes=1683
Combine input records=0
Combine output records=0
Reduce input groups=28090
Reduce shuffle bytes=16923444
Reduce input records=1647340
Reduce output records=28090
Spilled Records=3294680
Shuffled Maps =153
Failed Shuffles=0
Merged Map outputs=153
GC time elapsed (ms)=32323
CPU time spent (ms)=186490
Physical memory (bytes) snapshot=16614551552
Virtual memory (bytes) snapshot=97874968576
Total committed heap usage (bytes)=14691598336

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1087603226

File Output Format Counters

Bytes Written=957920