

```
2017-02-11 00:20:53,465 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-54-184.ec2.internal/172.31.54.184:8032
2017-02-11 00:20:57,164 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input paths to process : 10
2017-02-11 00:20:57,187 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2017-02-11 00:20:57,190 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev f7cb0596948c5bfd3e71d37b0f5bb21a19554666]
2017-02-11 00:20:57,607 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:10
2017-02-11 00:20:57,805 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1486772290921_0001
2017-02-11 00:20:58,830 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1486772290921_0001
2017-02-11 00:20:59,004 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-54-184.ec2.internal:20888/proxy/application_1486772290921_0001/
2017-02-11 00:20:59,005 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1486772290921_0001
2017-02-11 00:21:10,142 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486772290921_0001 running in uber mode : false
2017-02-11 00:21:10,144 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2017-02-11 00:21:27,269 INFO org.apache.hadoop.mapreduce.Job (main): map 10% reduce 0%
2017-02-11 00:21:28,277 INFO org.apache.hadoop.mapreduce.Job (main): map 17% reduce 0%
2017-02-11 00:21:29,289 INFO org.apache.hadoop.mapreduce.Job (main): map 27% reduce 0%
2017-02-11 00:21:30,298 INFO org.apache.hadoop.mapreduce.Job (main): map 64% reduce 0%
2017-02-11 00:21:31,303 INFO org.apache.hadoop.mapreduce.Job (main): map 87% reduce 0%
2017-02-11 00:21:32,309 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 0%
2017-02-11 00:21:33,313 INFO org.apache.hadoop.mapreduce.Job (main): map 97% reduce 0%
2017-02-11 00:21:36,330 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2017-02-11 00:21:43,366 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 9%
2017-02-11 00:21:44,371 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 11%
2017-02-11 00:21:45,374 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 21%
2017-02-11 00:21:46,379 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 61%
2017-02-11 00:21:47,383 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 63%
2017-02-11 00:21:48,387 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 89%
2017-02-11 00:21:51,398 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-02-11 00:21:52,414 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486772290921_0001 completed successfully
2017-02-11 00:21:52,686 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
    File System Counters
        FILE: Number of bytes read=27380199
        FILE: Number of bytes written=57028706
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=980
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=10
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=0
        S3: Number of bytes read=235553896
        S3: Number of bytes written=403049
        S3: Number of read operations=0
        S3: Number of large read operations=0
        S3: Number of write operations=0
    Job Counters
        Killed map tasks=1
        Killed reduce tasks=1
        Launched map tasks=11
        Launched reduce tasks=9
        Data-local map tasks=11
        Total time spent by all maps in occupied slots (ms)=9295968
        Total time spent by all reduces in occupied slots (ms)=12349248
        Total time spent by all map tasks (ms)=193666
```

```
Total time spent by all reduce tasks (ms)=128638
Total vcore-milliseconds taken by all map tasks=193666
Total vcore-milliseconds taken by all reduce tasks=128638
Total megabyte-milliseconds taken by all map tasks=297470976
Total megabyte-milliseconds taken by all reduce tasks=395175936
Map-Reduce Framework
  Map input records=7003910
  Map output records=7003910
  Map output bytes=259144670
  Map output materialized bytes=27227907
  Input split bytes=980
  Combine input records=0
  Combine output records=0
  Reduce input groups=2875
  Reduce shuffle bytes=27227907
  Reduce input records=7003910
  Reduce output records=2875
  Spilled Records=14007820
  Shuffled Maps =90
  Failed Shuffles=0
  Merged Map outputs=90
  GC time elapsed (ms)=6541
  CPU time spent (ms)=117130
  Physical memory (bytes) snapshot=10059681792
  Virtual memory (bytes) snapshot=74822901760
  Total committed heap usage (bytes)=9307160576
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=235553896
File Output Format Counters
  Bytes Written=403049
```