

```
2017-02-11 01:30:51,163 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to
ResourceManager at ip-172-31-57-40.ec2.internal/172.31.57.40:8032
2017-02-11 01:30:54,421 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input paths to process : 1
2017-02-11 01:30:54,431 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2017-02-11 01:30:54,433 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev f7cb0596948c5bfd3e71d37b0f5bb21a19554666]
2017-02-11 01:30:54,833 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:17
2017-02-11 01:30:54,987 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1486776526930_0001
2017-02-11 01:30:55,466 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1486776526930_0001
2017-02-11 01:30:55,558 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-57-40.ec2.internal:20888/proxy/application_1486776526930_0001/
2017-02-11 01:30:55,558 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1486776526930_0001
2017-02-11 01:31:06,811 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776526930_0001 running in uber mode : false
2017-02-11 01:31:06,813 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2017-02-11 01:31:24,986 INFO org.apache.hadoop.mapreduce.Job (main): map 1% reduce 0%
2017-02-11 01:31:28,003 INFO org.apache.hadoop.mapreduce.Job (main): map 2% reduce 0%
2017-02-11 01:31:31,024 INFO org.apache.hadoop.mapreduce.Job (main): map 4% reduce 0%
2017-02-11 01:31:32,027 INFO org.apache.hadoop.mapreduce.Job (main): map 5% reduce 0%
2017-02-11 01:31:34,044 INFO org.apache.hadoop.mapreduce.Job (main): map 10% reduce 0%
2017-02-11 01:31:35,047 INFO org.apache.hadoop.mapreduce.Job (main): map 13% reduce 0%
2017-02-11 01:31:36,055 INFO org.apache.hadoop.mapreduce.Job (main): map 16% reduce 0%
2017-02-11 01:31:37,060 INFO org.apache.hadoop.mapreduce.Job (main): map 21% reduce 0%
2017-02-11 01:31:38,064 INFO org.apache.hadoop.mapreduce.Job (main): map 28% reduce 0%
2017-02-11 01:31:39,070 INFO org.apache.hadoop.mapreduce.Job (main): map 32% reduce 0%
2017-02-11 01:31:40,076 INFO org.apache.hadoop.mapreduce.Job (main): map 39% reduce 0%
2017-02-11 01:31:41,079 INFO org.apache.hadoop.mapreduce.Job (main): map 43% reduce 0%
2017-02-11 01:31:42,083 INFO org.apache.hadoop.mapreduce.Job (main): map 48% reduce 0%
2017-02-11 01:31:43,087 INFO org.apache.hadoop.mapreduce.Job (main): map 52% reduce 0%
2017-02-11 01:31:44,091 INFO org.apache.hadoop.mapreduce.Job (main): map 56% reduce 0%
2017-02-11 01:31:45,096 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 0%
2017-02-11 01:31:46,102 INFO org.apache.hadoop.mapreduce.Job (main): map 65% reduce 0%
2017-02-11 01:31:47,106 INFO org.apache.hadoop.mapreduce.Job (main): map 66% reduce 0%
2017-02-11 01:31:48,109 INFO org.apache.hadoop.mapreduce.Job (main): map 67% reduce 0%
2017-02-11 01:31:49,113 INFO org.apache.hadoop.mapreduce.Job (main): map 68% reduce 0%
2017-02-11 01:31:52,123 INFO org.apache.hadoop.mapreduce.Job (main): map 69% reduce 0%
2017-02-11 01:31:53,126 INFO org.apache.hadoop.mapreduce.Job (main): map 73% reduce 0%
2017-02-11 01:31:54,131 INFO org.apache.hadoop.mapreduce.Job (main): map 81% reduce 0%
2017-02-11 01:31:55,134 INFO org.apache.hadoop.mapreduce.Job (main): map 89% reduce 0%
2017-02-11 01:31:56,138 INFO org.apache.hadoop.mapreduce.Job (main): map 94% reduce 0%
2017-02-11 01:31:58,145 INFO org.apache.hadoop.mapreduce.Job (main): map 96% reduce 0%
2017-02-11 01:31:59,148 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 0%
2017-02-11 01:32:00,153 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2017-02-11 01:32:07,177 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 7%
2017-02-11 01:32:08,180 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 23%
2017-02-11 01:32:10,188 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 43%
2017-02-11 01:32:11,196 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 60%
2017-02-11 01:32:12,200 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 65%
2017-02-11 01:32:13,204 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 77%
2017-02-11 01:32:14,207 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 83%
2017-02-11 01:32:15,214 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 94%
2017-02-11 01:32:16,220 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-02-11 01:32:17,230 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1486776526930_0001 completed successfully
2017-02-11 01:32:17,339 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
File System Counters
FILE: Number of bytes read=140015225
FILE: Number of bytes written=297677204
```

FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=1683  
HDFS: Number of bytes written=0  
HDFS: Number of read operations=17  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=0  
S3: Number of bytes read=1087575223  
S3: Number of bytes written=975899  
S3: Number of read operations=0  
S3: Number of large read operations=0  
S3: Number of write operations=0

#### Job Counters

Killed map tasks=2  
Killed reduce tasks=1  
Launched map tasks=18  
Launched reduce tasks=10  
Data-local map tasks=18  
Total time spent by all maps in occupied slots (ms)=35802528  
Total time spent by all reduces in occupied slots (ms)=15146880  
Total time spent by all map tasks (ms)=745886  
Total time spent by all reduce tasks (ms)=157780  
Total vcore-milliseconds taken by all map tasks=745886  
Total vcore-milliseconds taken by all reduce tasks=157780  
Total megabyte-milliseconds taken by all map tasks=1145680896  
Total megabyte-milliseconds taken by all reduce tasks=484700160

#### Map-Reduce Framework

Map input records=30868726  
Map output records=30868726  
Map output bytes=895193054  
Map output materialized bytes=154366719  
Input split bytes=1683  
Combine input records=0  
Combine output records=0  
Reduce input groups=28980  
Reduce shuffle bytes=154366719  
Reduce input records=30868726  
Reduce output records=28980  
Spilled Records=61737452  
Shuffled Maps =153  
Failed Shuffles=0  
Merged Map outputs=153  
GC time elapsed (ms)=15718  
CPU time spent (ms)=267100  
Physical memory (bytes) snapshot=17481207808  
Virtual memory (bytes) snapshot=97884065792  
Total committed heap usage (bytes)=16482566144

#### Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=1087575223

#### File Output Format Counters

Bytes Written=975899